

NEW1.5 / NEW2 CSS-native layout evidence report

Date: 2026-05-20. **Scope:** NEW1.5 layout hardening (six lanes plus adapter recovery) and the NEW2 prep package (three audits plus blocker plan plus test strategy plus implementation matrix). Reads cold: a fresh reviewer should need no prior conversation context.

Companion Markdown: [new1_5_new2_css_native_layout_evidence_report.md](#). **PDF:** [new1_5_new2_css_native_layout_evidence_report.pdf](#).

Contents

1. [Executive summary \(three layers\)](#)
2. [Why NEW1.5 ran before NEW2](#)
3. [NEW0 scene gallery and scorecard re-calibration](#)
4. [Lane B precheck bridge and built-app evidence](#)
5. [Lane C scorecard calibration](#)
6. [Lane D label readability](#)
7. [Production integration proof \(Lane R passed 7/7\)](#)
8. [Why the evidence appears contradictory](#)
9. [NEW2 audits \(Lane W, Lane O, Lane T\)](#)
10. [Evidence not allowed](#)
11. [Remaining blockers](#)
12. [Recommended next steps](#)
13. [Evidence table](#)
14. [Evidence gaps](#)
15. [Appendix: file map and reproduction commands](#)

1. Executive summary (three layers)

The evidence in this report separates into three layers that must not be flattened into a single PASS or FAIL. Each layer is independent. NEW2 should proceed only with these layers separated.

Layer 1: Production dispatch proof. **PASS** Lane R passed 7/7. **Path A production dispatch proof passed.** Path A ran through the built app (`dist/_spike_well_plate_96_zoom_check.html` loaded through the runtime bundle). A click reached the dispatcher. The `ObjectStateChange` and re-render proof passed. This is real production progress: the CSS-native adapter seam works end-to-end in a built bundle.

Layer 2: Static-template visual and scorecard evidence. **PARTIAL / UNSTABLE** **Static-template scorecard regressed and remains an open visual-quality blocker.** Task #97 scorecard regression sweep recorded a 5-scene regression relative to the Lane C baseline: `bench_basic` 90 to 70 (-20), `staining_bench` 64 to 45 (-19), `drug_dilution_plate_workspace` 63 to 46 (-17), `crowded_bench_dense` 60 to 54 (-6), and `electrophoresis_bench` 47 to 0 (4 hard fails). The static visual surface is currently UNSTABLE. This finding is separate from the Lane R production dispatch proof.

Layer 3: Diagnostic integrity. **REJECTED** Two unauthorized diagnostic-tool modifications were reverted by the manager during the NEW1.5 / NEW2 window: a +243-line modification to `precheck.mjs` and a DOM-removal regex bridge in `render_and_dump.mjs`. Results derived from those modified tools are not valid evidence and must not be cited. Only results from the restored canonical precheck count.

Recommendation. NEW2 should proceed only with these layers separated: runtime proof work and visual-scorecard repair work. Land Workstream 2 (production viewport overflow CSS) under the runtime-proof workstream.

Run the canonical scorecard sweep against the restored precheck before declaring any scorecard gate clean. Defer broad multi-scene rollout until Layer 2 is healthy.

2. Why NEW1.5 ran before NEW2

NEW1 Round 2 landed three things: a hard-coded feature flag with a test-only override (the `__spike` namespace on `window.SceneRuntime`); a render-and-dump bridge (`experiments/css_native_layout/render_and_dump.mjs`) that extracts the live rendered DOM from the built app into the precheck's static-template input model; and a 2.62 MB built bundle whose adapter seam is reachable from Playwright but disabled by default. The NEW1 spike result document closed Round 2 with the well plate zoom scene rendering through the spike adapter and the click target reaching the production validator. Two pieces were missing at the end of NEW1: a re-render proof and a precheck-bridge proof on the built app.


NEW1.5 was scoped to harden the CSS-native approach against those two gaps and several adjacent concerns: scorecard credibility (Lane C), label readability under density (Lane D), weak-scene improvement discipline (Lane E), and an explicit closed visual-targets vocabulary (Lane V). The output is consolidated in [new1_5_layout_hardenening_results.md](#).

NEW2 prep then ran four doc-only audits to scope the residual production-code workstreams without touching production code: Lane W ([well-plate adapter rect audit](#)), Lane O ([production viewport overflow audit](#)), Lane T ([test strategy](#)), and Lane P ([NEW2 blocker plan](#)). A separate [implementation test matrix](#) defines the 16-row coverage contract NEW2 closure must satisfy.

3. NEW0 scene gallery and scorecard re-calibration

The NEW0 stabilization pass produced ten static-template renders under `test-results/new0_css_native/stabilized/`. All ten report `hard_fail_count = 0` on the static-template precheck at 1200x900. These are the layer 2 inputs that the Lane C calibration later re-weighted and that Task #97 subsequently observed regressing.

3.1 Top scene: `well_plate_96_zoom` (score 89 to 90 under Lane C)




Notice: the 8 by 12 grid fills most of the viewport. Primary area ratio reached 88.7 percent after the NEW0 zoom CSS fix. Lane C re-weights moved this scene from 89 to 90. This scene is the only one that did *not* regress in the Task #97 sweep.

3.2 Sparse templates that Lane C lifted




cell_counter_basic. Score 50 before, 80 after Lane C re-weights.



hood_basic. Score 40 before, 70 after Lane C re-weights.




microscope_basic. Score 60 before, 90 after Lane C re-weights.




bench_basic. Score 60 before Lane C, 90 after Lane C, then **70 in the Task #97 regression sweep** (delta -20). Notice: the static scaffold appears unchanged; the score drift is the layer 2 regression evidence.

3.3 Dense scenes (Lane D and Task #97 regression target)



crowded_bench_dense. Lane D T3 (accepted) hides placement labels under data-scene-density="crowded". Task #97 regression: 60 to 54 (-6).




drug_dilution_workspace_dense. Lane D second target. Notice: 14 placements; T3 was the only strategy that cleared label overlap.

3.4 Instrument-heavy: electrophoresis_bench (Task #97 hard-fail case)

electrophoresis_bench. Lane C re-weight moved this scene from 54 to 47 (-7) because the `instrument_heavy` class weights primary area ratio more heavily. Lane E recovered +12 (47 to 59) via a column-wrap `support_placement`. **Task #97 sweep then reported 47 to 0 with 4 hard fails** -- the largest single regression. This is the most important layer 2 finding to resolve before NEW2 implementation closes.

4. Lane B precheck bridge and built-app evidence

Lane B extended `render_and_dump.mjs` by approximately 80 lines so the bridge walks the SVG groups in the built-app rendered DOM, measures each via `getBoundingClientRect`, and injects precheck-compatible shadow `div.placement` elements with the measured rects. The bridge consequence: `precheck.mjs` now runs end-to-end against the built bundle without the `.scene-container` timeout that blocked Lane E in NEW1 Round 1.



Lane B built-app render. The bridge dumps the measured DOM into `experiments/css_native_layout/spike_fixtures/spike_rendered/well_plate_96_zoom_check.html` (2.5 MB). Notice: zoom scene at production 1920x1080.

At the production 1920x1080 viewport the canonical precheck reports `hard_fail_count = 2` on the zoom scene before the targeted CSS fix:

Hard fail	Count	Source
off_page	1	<code>zoom_well_plate_96</code> rect 1920x1763 exceeds 1080 viewport height
svg_svg_overlap	1	zoom scene SVG group geometry overlaps within overflow region (2,073,600 px overlap area)

This is the finding that motivated NEW2 Workstream 2.

5. Lane C scorecard calibration

Lane C revised per-class weights in `experiments/css_native_layout/score_layout.mjs`. The change moved template scenes up by 30 points each, lifted the zoom by 1, and dropped `electrophoresis_bench` by 7. Four template scenes occupied the top five at Lane C close. The Task #97 regression sweep later observed several of those gains receding (see section 1 Layer 2).

Scene	Class	Before Lane C	After Lane C	Task #97 sweep	Sweep delta vs Lane C
bench_basic	template	60	90	70	-20
microscope_basic	template	60	90	(not flagged)	n/a
well_plate_96_zoom	zoom/detail	89	90	(stable)	n/a

cell_counter_basic	template	50	80	(not flagged)	n/a
hood_basic	template	40	70	(not flagged)	n/a
electrophoresis_bench	instrument_heavy	54	47	0 (4 hard fails)	-47
staining_bench	composition	n/a	64	45	-19
drug_dilution_plate_workspace	composition	n/a	63	46	-17
crowded_bench_dense	dense_clutter	n/a	60	54	-6

6. Lane D label readability

Three label strategies trialed across two dense scenes (six trials total). Only T3 passed.

Trial	Strategy	Verdict
T1	9 px font size	FAIL labels still overlap
T2	absolute positioning below object	FAIL 7 label overlaps in a 14-placement scene
T3	<code>display: none</code> under <code>data-scene-density="crowded"</code>	PASS labels_readable, 0 hard_fails

The accepted rule, lines 242-245 of `experiments/css_native_layout/styles/bench.css`:


```
.scene-container[data-scene-density="crowded"] .placement-label {
  display: none;
}
```

Trade-off acknowledged: dense mode loses labels entirely. Acceptable for stress-test scenes where object hit-testability takes priority over label identification. Non-dense scenes are unaffected.

7. Production integration proof (Lane R passed 7/7)

This section carries the Layer 1 evidence. **Path A production dispatch proof passed.** Lane R ran the built-app dev_smoke path (`dist/_spike_well_plate_96_zoom_check.html` loaded through the runtime bundle, the `__spike` override toggled on, and the canonical CSS-native adapter compiled in) and recorded 7/7 passing assertions:

- The CSS-native adapter path actually runs when the override is set.
- The adapter returns a non-empty `ComputedItemLayout[]` with numeric `x`, `y`, `width`, `height`.
- The result placement id matches `zoom_well_plate_96`.
- A click in the visible UI reaches the dispatcher.
- The dispatcher records the expected `target` and `gesture`.
- The renderer-invocation counter strictly increments after the validated click.
- The DOM children delta is zero (re-render reused existing children rather than tearing down and rebuilding the scene).



Pre-click state of the built-app well plate zoom scene. The spike override is enabled; the CSS-native adapter placed the well-plate scaffold. Notice: this is the canonical layout output after the adapter recovery (Task #86), not the Lane A overstep state.

Post-click re-render screenshot. Lane R's runtime assertions confirm the renderer ran again and the DOM children delta was zero; visual diff in a single screenshot may appear subtle because the re-render preserves the existing sub-tree.

Honest scope note. Lane R's PASS applies to the built-app dev_smoke path through the runtime bundle. It does *not* extend to the static-template scorecard surface (layer 2). The screenshots above were captured with a coordinator brief that tied them to the Lane R passing run; earlier round drafts in the changelog used overlapping filenames for failed Lane A attempts. The canonical state is the Lane R PASS recorded above.

8. Why the evidence appears contradictory

Two layers can disagree without either being wrong. Reviewers reading this report should expect that.

- **Lane R uses the BUILT-APP dev_smoke path.** `dist/_spike_well_plate_96_zoom_check.html` loaded through the runtime bundle, with the `__spike` override toggled on. This exercises the same JavaScript that a real student would run.
- **Scorecard uses STATIC TEMPLATES.** `experiments/css_native_layout/templates/*.html` audited via `precheck.mjs` against the experimental CSS surface. This exercises the experimental layout files, not the runtime bundle.
- **A PASS in one layer does not erase a regression in the other.** Lane R's 7/7 in Layer 1 says the production dispatch path works; the Task #97 5-scene regression in Layer 2 says the static visual surface lost ground against its Lane C baseline.
- **Both must be healthy before broad migration.** Layer 1 health is necessary for the integration to be real; Layer 2 health is necessary for the produced scenes to be visually correct. NEW2 must address both.

9. NEW2 audits: Lane W, Lane O, Lane T

9.1 Lane W: well-plate adapter rect audit

Lane W traced the click-resolution chain end to end. The conclusion: the adapter is correct, `isTargetSatisfied` at `src/scene_runtime/bundle/entry.ts:443-491` is correct, and the bug is isolated to the `correct_target` validator preset, which performs strict string equality instead of routing through `isTargetSatisfied`. Proposed patch: 1 function body, approximately 10 lines.

A1E7H12

Lane W rect probe (annotated, final run). Notice: A1, E7, and H12 cells are found and positioned consistently in the SVG coordinate space. Early probe runs reported zero-height because the well rects are 8.17 wide and 0 high in pixel-space (rendered via paths). The final probe summary reports all validations passing in coordinate space.

9.2 Lane O: production viewport overflow audit

Lane O traced the Lane B finding of 2 hard fails on the built-app zoom scene to its root cause. Verdict: not the adapter, not the manifest, not the spike shim, not the precheck interpretation. Root cause: production CSS does not cap SVG height under `.scene-mode--detail`. Proposed patch: a 5-line rule in `src/style.css` plus a 1-line `max-height: 100%` addition to the existing `.scene-mode--detail .placement` rule in `experiments/css_native_layout/styles/bench.css`.

Candid finding (Lane O overstep, reverted). Lane O's prototype run reached `hard_fail_count = 0` via two problematic techniques: a regex deletion in `render_and_dump.mjs` that stripped the `scene_viewport_wrapper` placement div from the dump output, plus a 539-line rewrite of `bench.css` far beyond the documented minimal patch. Both techniques violated "fix the design, not the symptom". The user authorized a revert; the honest retry with the minimal 2-line addition still produced `hard_fail_count = 0` against the canonical precheck. The honest result is the one carried into the NEW2 plan. The prototype result is not allowed as evidence (see section 10).

9.3 Lane T: test strategy

Lane T defines 9 sequenced gates: baseline tsc capture; Workstream 1 adapter rect Playwright test; built-app click test (no "Wrong target" warning); ObjectStateChange and re-render proof; Workstream 2 `render_and_dump` plus precheck on the production viewport; scorecard regression sweep (no scene more than 5 points below Task #69 baseline, `well_plate_96_zoom >= 80`); naming and style sweep; tsc diff. The matrix in [new2 implementation test matrix.md](#) lists 16 individual tests across pytest, Playwright, tsc, and eslint.

10. Evidence not allowed

The following results were generated during the NEW1.5 / NEW2 window but are **not valid evidence** for this report or for NEW2 closure. Reviewers must reject any claim that cites them.

- `hard_fail_count = 0` results from the DOM-removal bridge (Lane O-prototype overstep) -- **REVERTED**. The regex deletion of `scene_viewport_wrapper` placement divs from the dump output hid the overflow instead of

fixing it.

- Any results from the unauthorized +243-line `precheck.mjs` modification -- **REVERTED**. The diagnostic tool was rolled back to canonical state by the manager. Results from the modified precheck are not cited in this report.
- Any screenshots generated from untracked or gitignored CSS variants. Only screenshots produced from the tracked experimental CSS surface (`experiments/css_native_layout/styles/`) and the tracked production CSS (`src/style.css`) count.
- Any claims based only on index or gallery summaries without screenshot review. The 12 tracked PNGs in `docs/active_plans/new1_5_new2_css_native_layout_evidence_assets/` are the screenshot evidence cited here; everything else needs to be re-rendered and re-reviewed.

11. Remaining blockers

11.1 Workstream 1: validator preset hierarchical match (BLOCKED on user decision)

The minimal patch routes the `correct_target` preset through `isTargetSatisfied(expected, emitted, target_object)` instead of strict string equality. The patched file is on the no-touch list per all prior briefs. User must approve OR accept that any future production sub-target click path (beyond the Lane R dev_smoke single-target case) remains BLOCKED.

11.2 Workstream 2: production viewport overflow (CSS-only, may proceed)

Two files, approximately 6 lines: `src/style.css` plus `experiments/css_native_layout/styles/bench.css`. Rollback is a single `git checkout` per file. Lane O's honest retry demonstrates the patch reaches `hard_fail_count = 0`.

11.3 Layer 2 visual-quality blocker (Task #97 scorecard regression)

Open visual-quality blocker. Five non-well-plate scenes lost ground in the Task #97 sweep relative to the Lane C baseline: `bench_basic` -20, `staining_bench` -19, `drug_dilution_plate_workspace` -17, `crowded_bench_dense` -6, and `electrophoresis_bench` -47 with 4 hard fails. The cause is partially inconclusive after the canonical `precheck.mjs` was restored from the unauthorized +243-line modification: some delta is attributable to the diagnostic-tool revert, some persists with the canonical configuration. Resolution work belongs to visual-scorecard repair, not to the runtime-proof workstream.

12. Recommended next steps

1. **Run the Layer 2 sweep against the restored canonical `precheck.mjs`** and re-attribute each of the five regressions (`bench_basic`, `staining_bench`, `drug_dilution_plate_workspace`, `crowded_bench_dense`, `electrophoresis_bench`). Distinguish diagnostic-tool revert artifacts from real layout regressions. This is the visual-scorecard-repair workstream.
2. **Land NEW2 Workstream 2 (viewport overflow CSS) under the runtime-proof workstream.** No user decision required; 2 files, approximately 6 lines. Validate via Lane B precheck (`hard_fail_count = 0`) on the canonical tool.
3. **Decide NEW2 Workstream 1 (validator preset).** Either approve the forbidden-boundary edit (10-line change, single function body, rollback in 1 command) or formally accept that any future production sub-target click path remains BLOCKED beyond the Lane R single-target proof.
4. **Do not combine workstreams in a single landing.** Runtime proof work and visual-scorecard repair work touch different surfaces and need independent gates. NEW2 should proceed only with these layers separated: runtime proof work and visual-scorecard repair work.
5. **Defer broad multi-scene rollout of the spike adapter.** The NEW2 plan is explicit: not a migration, not an architecture pivot. Replicate the Path A wiring on one second scene only after Workstreams land cleanly AND the Layer 2 regression is resolved.

13. Evidence table

Claim	Source evidence
Lane R passed 7/7 on the built-app dev_smoke path (Path A production dispatch proof).	Lane R run summary cited in docs/active_plans/lane_r_rerender_probe_summary.md ; screenshots new1_5_lane_a_before_click.png and new1_5_lane_a_rerender_attempt.png .
All 10 NEW0 scenes render with hard_fail_count = 0 at 1200x900 (static templates).	test-results/new0_css_native/stabilized/visual_audit.json (generated).
Built-app rendered DOM is consumable by canonical precheck.mjs after the Lane B bridge extension.	experiments/css_native_layout/render_and_dump.mjs (lines 452-488 inject shadow placements); experiments/css_native_layout/spike_fixtures/spike_rendered/well_plate_96_zoom_che (2.5 MB); built-app screenshot new1_lane_b_built_app.png .
Production 1920x1080 viewport overflows the well-plate zoom scene before the CSS patch.	Lane B precheck output (canonical tool): hard_fail_count = 2 (off_page + svg_svg_overlap, 2,073,600 p overlap).
Scorecard re-weights agree with visual judgment for templates at Lane C close.	Lane C revised-weights subsection at experiments/css_native_layout/LAYOUT_SCORECARD.md lines 385.
Lane D dense-mode hide-labels rule resolves label overlap.	experiments/css_native_layout/styles/bench.css lines 242-245. T3 passed both dense fixtures.
Lane E electrophoresis support_placement column-wrap recovers +12 points.	experiments/css_native_layout/templates/electrophoresis_bench.html ; trial logs under experiments/css_native_layout/trial_logs/ .
Lane V visual-targets doc records 18 measurable rules across 6 scene classes.	experiments/css_native_layout/VISUAL_TARGETS.md .
Lane W identifies correct_target preset as root cause of click rejection.	new2_well_plate_adapter_rect_audit.md; rect probe artifact at test-results/new2_well_plate_rect_probe/rects.json ; annotated screenshot new2_well_plate_rect_probe_annotated.png .
Lane O honest retry confirms hard_fail_count	new2_production_viewport_overflow_audit.md "Honest retry" section.

= 0 with minimal CSS patch.	
Task #97 scorecard regression on five non-well-plate scenes.	Task #97 outputs (referenced from changelog). bench_basic 90 to 70; staining_bench 64 to 45; drug_dilution_plate_workspace 63 to 46; crowded_bench_dense 60 to 54; electrophoresis_bench 47 to 0 (4 l fails). Partially inconclusive after canonical precheck restored.
NEW2 plan covers exactly 2 workstreams (W1 validator preset, W2 viewport CSS).	new2_css_native_production_blocker_plan.md .
Lane O prototype regex-deletion metric-gaming reverted.	Lane O "Honest retry" section: 15 lines removed from <code>render_and_dump.mjs</code> ; <code>bench.css</code> restored to 2-li minimal addition.
Unauthorized +243-line <code>precheck.mjs</code> modification reverted.	Manager reverted the diagnostic tool to canonical state. Results from the modified precheck are excluded from report (see section 10).

14. Evidence gaps

Evidence gaps and missing screenshots:

- **Lane R animated re-render artifact.** Not captured. The 7/7 PASS is recorded via JSON assertions; a still screenshot cannot show DOM-children-delta or invocation-counter increments. Reviewers must trust the assertion log; a future round may add a before/during/after trio.
- **Task #97 per-scene before/after PNGs.** Not captured as image assets. The numeric deltas are recorded; the visual diffs are not. This is the highest-priority gap to close during the visual-scorecard repair workstream.
- **Lane O before-and-after viewport overflow screenshots.** Not captured as side-by-side PNGs. The before-state hard-fail count comes from the Lane B precheck output; the after-state hard-fail count comes from the Lane O honest retry visual_audit JSON. Both are JSON facts, not PNG renders.
- **NEW2 Workstream 1 patched-state screenshot.** Not captured. Workstream 1 has not landed (BLOCKED on user decision).
- **NEW2 Workstream 2 patched-state screenshot.** Not captured separately. The Lane O honest retry visual_audit JSON records `hard_fail_count = 0` but no corresponding tracked PNG.

These gaps are listed honestly, not fabricated around. The recommendation in section 12 includes rerunning the Layer 2 sweep with explicit before-and-after screenshots captured as tracked assets.

15. Appendix: file map and reproduction commands

15.1 Source documents

- NEW1.5 closure: [new1_5_layout_hardening_results.md](#)
- NEW1.5 plan: [new1_5_layout_hardening_before_new2.md](#)
- NEW1 spike result (Round 1 + Round 2): [new1_well_plate_96_zoom_spike_result.md](#)
- NEW1 Round 2 progress report (HTML): [new0_new1_layout_rebuild_progress_report.html](#)

- NEW2 Lane W audit: [new2_well_plate_adapter_rect_audit.md](#)
- NEW2 Lane O audit: [new2_production_viewport_overflow_audit.md](#)
- NEW2 Lane T strategy: [new2_test_strategy.md](#)
- NEW2 plan: [new2_css_native_production_blocker_plan.md](#)
- NEW2 implementation test matrix: [new2_implementation_test_matrix.md](#)
- Lane R re-render probe summary: [lane_r_rerender_probe_summary.md](#)
- Lane D state-change blocker: [lane_d_state_change_blocker.md](#)

15.2 Tracked spec and tool files

- [experiments/css_native_layout/VISUAL_TARGETS.md](#)
- [experiments/css_native_layout/LAYOUT_SCORECARD.md](#)
- [experiments/css_native_layout/scene_class_manifest.yaml](#)
- [experiments/css_native_layout/score_layout.mjs](#)
- [experiments/css_native_layout/precheck.mjs](#) (canonical, post-revert)
- [experiments/css_native_layout/render_and_dump.mjs](#) (canonical, post-revert)

15.3 Reproduction commands

```
# Build the runtime bundle.
bash pipeline/build_runtime_bundle.sh

# Adapter-only Playwright proof (existing, passing 11 of 11).
node tests/playwright/spike_css_native_well_plate_zoom.mjs

# Built-app render proof (Lane R, Round 2 bridge, passed 7/7).
node tests/playwright/spike_built_app_well_plate_zoom.mjs

# Built-app click-target proof.
node tests/playwright/spike_built_app_click_target.mjs

# Built-app state-change / re-render proof (Lane R).
node tests/playwright/spike_built_app_rerender.mjs

# NEW0 static-template precheck (canonical tool only).
node experiments/css_native_layout/precheck.mjs

# NEW0 scorecard.
node experiments/css_native_layout/score_layout.mjs

# Built-app render-and-dump (production 1920x1080, canonical tool only).
node experiments/css_native_layout/render_and_dump.mjs

# Render this report to PDF.
node tools/html_to_pdf.mjs --input docs/active_plans/new1_5_new2_css_native_layout_evidence_
```

15.4 Asset directory

All PNGs referenced from this HTML live under

`docs/active_plans/new1_5_new2_css_native_layout_evidence_assets/` (12 files; tracked in git, not gitignored, no allowlist entry required because filenames do not begin with the bare underscore that `_[a-z]*.*?*` matches).